

# Let your Fabric work!

## Load testing Microsoft Fabric

Reitse Eskens

# Definitions

# A large amount of data



# Load testing



# Reitse Eskens

Engineer | Architect

Axians Business Analytics



Reitse Eskens@axians.com

/in/reitseeskens

<https://sqlreitse.com>

@2meterDBA



# Coming from a SQL world



# From definitions to Monitoring



# Learning a new tool: Fabric Capacity Metrics



# About capacity units

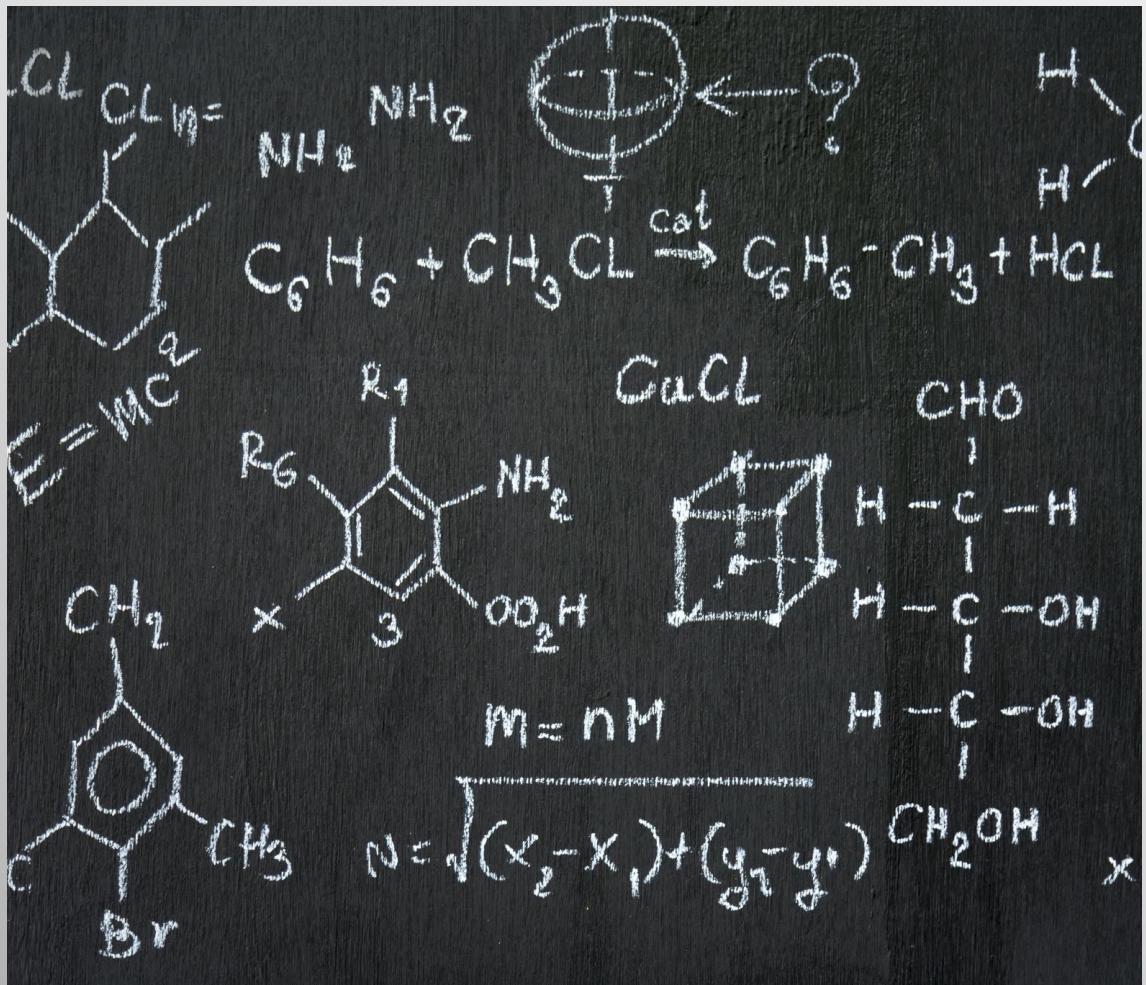
1 second of work = 1 capacity unit

F2 equals to 2 capacity units per second

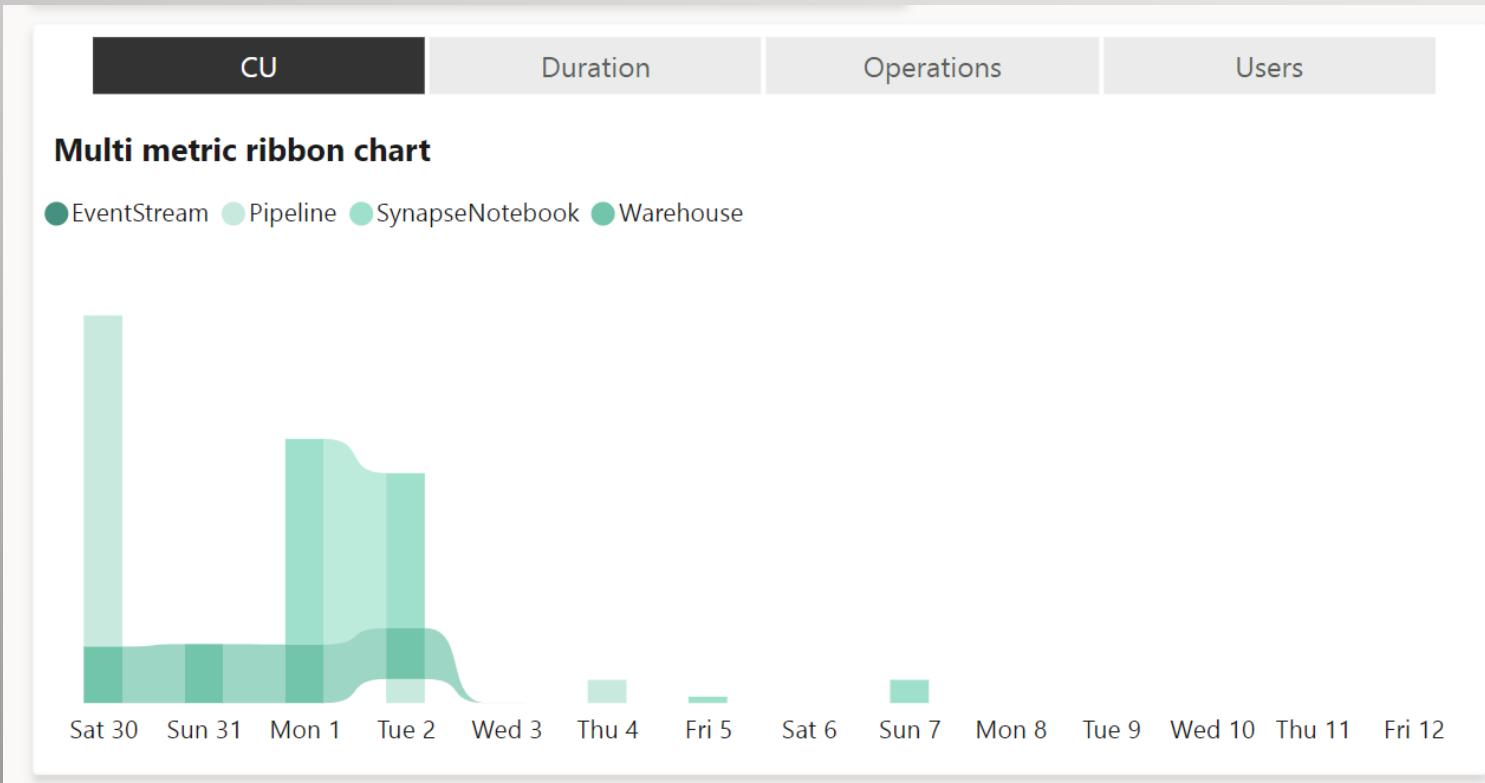
Capacity units are evaluated in blocks of 30 seconds

$F_2 = 2 * 30s = 60$  capacity units

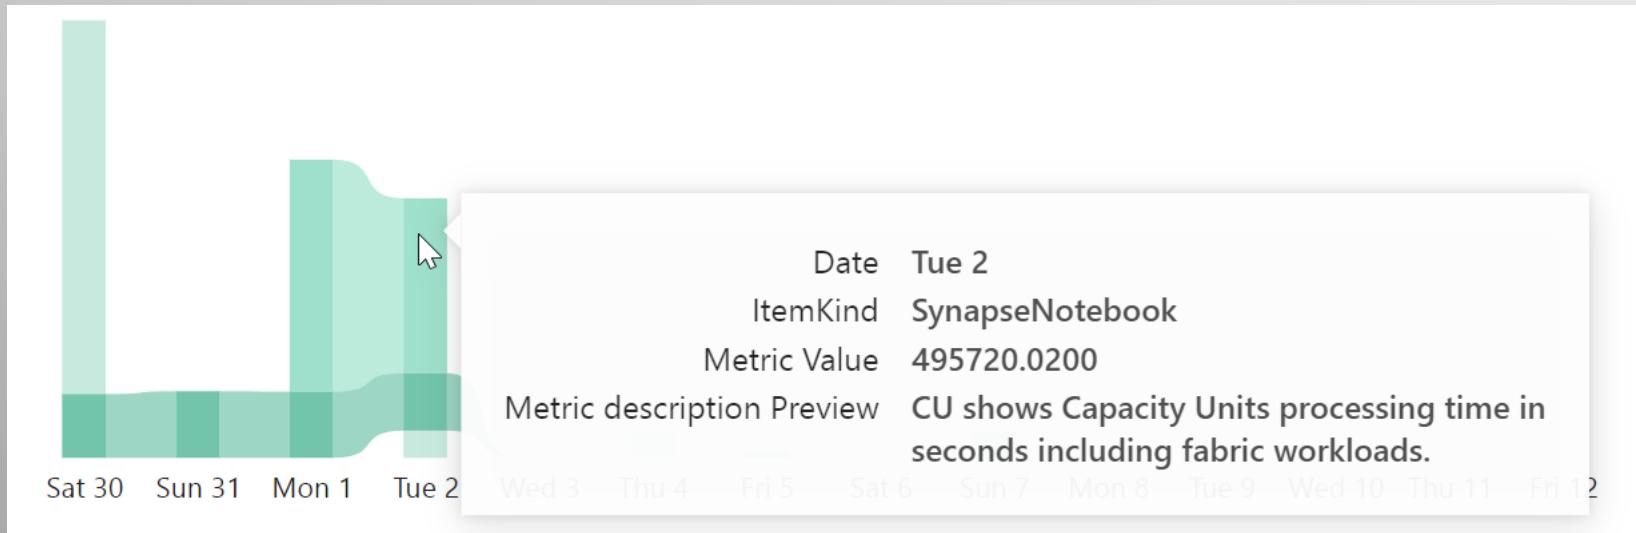
$F_{64} = 64 * 30s = 1920$  capacity units



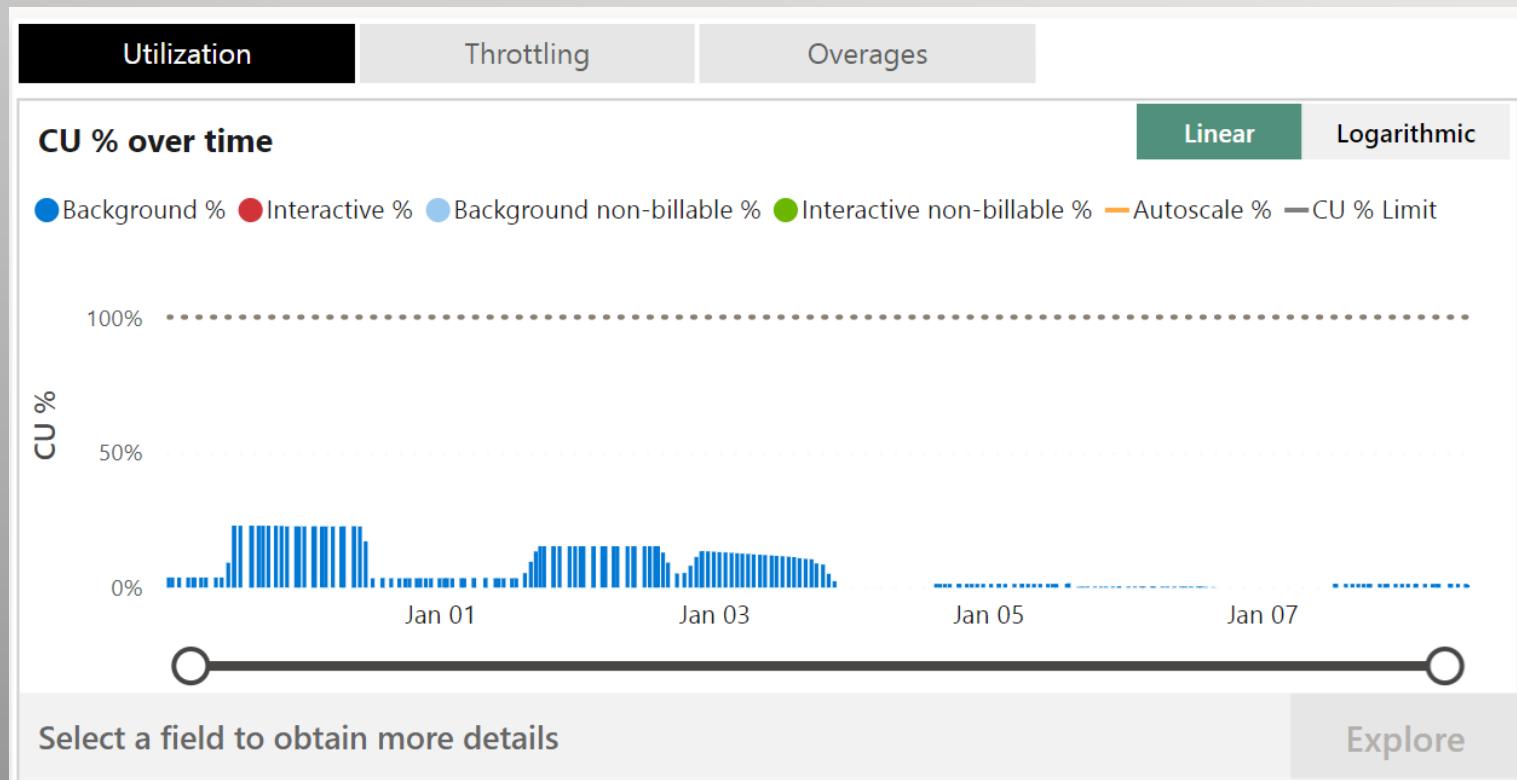
# Fabric Capacity Metrics: Multi metric ribbon chart



# Fabric Capacity Metrics: Multi metric ribbon chart



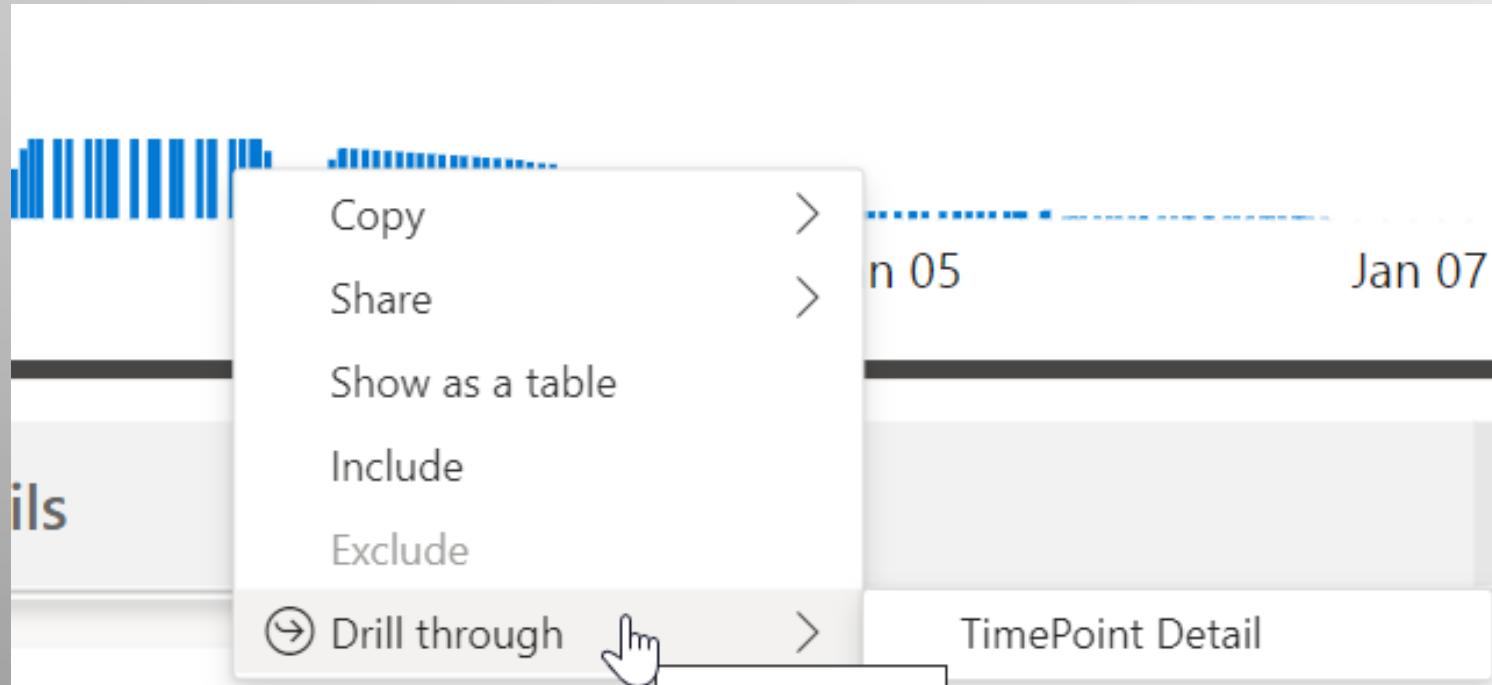
# Fabric Capacity Metrics: Capacity Utilization and Throttling



# Fabric Capacity Metrics: Capacity Utilization and Throttling

Policy	Consumption	Impact
Overage protection	Usage $\leq$ 10 minutes	Jobs can consume 10 minutes of future capacity use without throttling.
Interactive delay	10 minutes $<$ usage $\leq$ 60 minutes	User requested interactive jobs are throttled.
Interactive rejection	60 minutes $<$ usage $\leq$ 24 hours	User requested interactive jobs are rejected.
Background rejection	Usage $>$ 24 hours	User scheduled background jobs are rejected and not executed.

# Fabric Capacity Metrics: Drill through



- Large amounts of data is a matter of perspective
- Microsoft Fabric allows you to learn new skills
- Measuring Fabric capacity usage is one of them

# From theory to data

$\text{O. } \triangle \because x^2 + y^2 = ab + 4c$   $\text{Circ } c(x, y) \left\{ \begin{array}{l} xy = c \\ cx - cy = ab^2 \\ \pi = c \end{array} \right.$   $\text{AB}$

$\boxed{\beta = 70^\circ} \quad A = B$   $A \cap B, \frac{24+8}{Y} + \frac{2^2+3^2}{C} + \overrightarrow{x} \rightarrow g$

$\text{men} = 584. + n^{30} (x^2 + 34x + c)$

$x = 9.23$   $\sum_{x=2}^{n=14!} N_{50} \cdot x - \frac{1}{2} [984 + x^2 + \rho_{ab}]$

$\rightarrow x \leq 549$

## Preparing the data

- TPC-H
- .\dbgen.exe -s 10 -C 3 -v -S 1
  - -s 10 = 10 GB of data
  - -C 3 = three sets of tables
  - -v = verbose
  - -S 1 = first set of the three sets



## Preparing the data

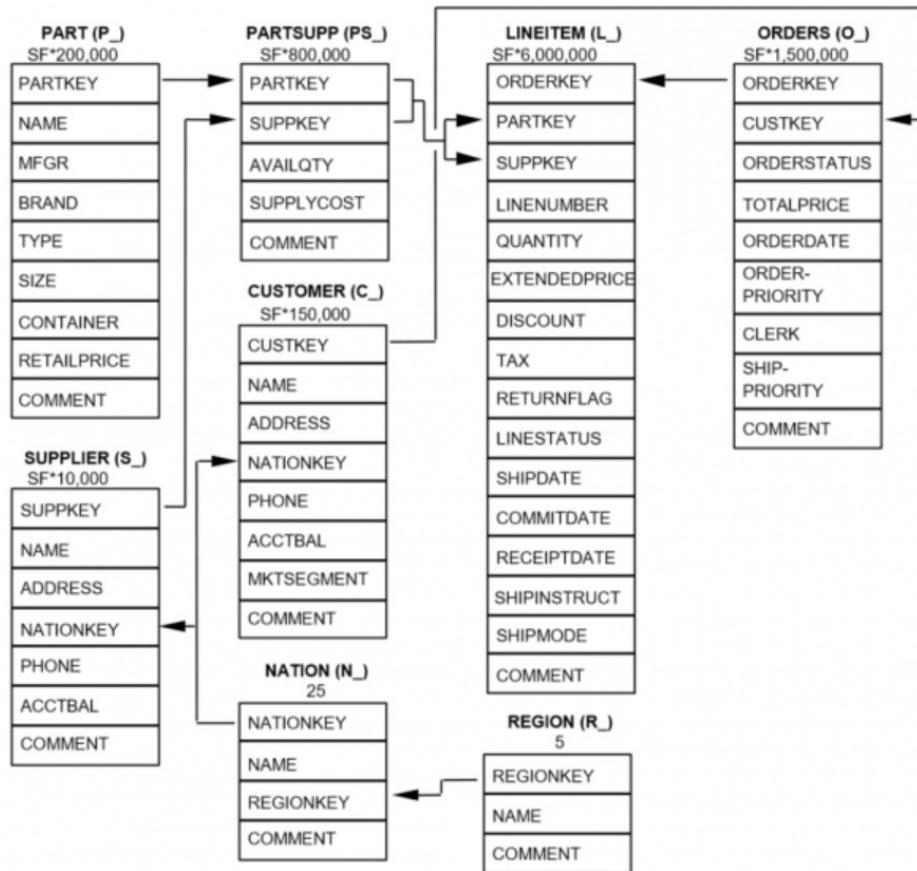
- .\dbgen.exe -v -U200
  - -v = verbose
  - -U 200 = 200 sets of update and delete sets



# Database schema

## 1 Database schema

The TPC-H Schema

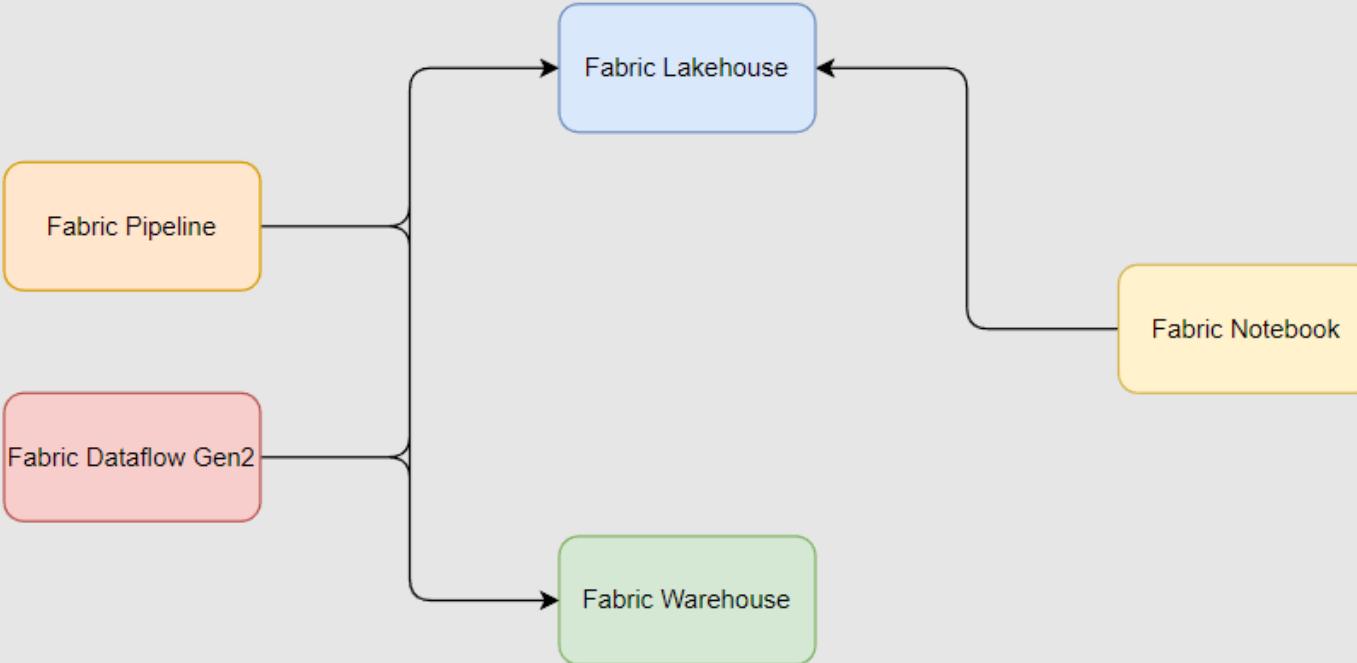


# Files

Name	Access Tier	Access Tier Last Modified	Last Modified	Blob Type	Content Type	Size	Status
region.tbl.csv	Hot (inferred)		2-1-2024 15:45	Block Blob	text/plain	394 B	Active
<a href="#">nation.csv</a>	Hot (inferred)		2-1-2024 15:45	Block Blob	text/plain	2,20 KiB	Active
supplier.tbl.1.csv	Hot (inferred)		2-1-2024 15:08	Block Blob	text/plain	137,20 MiB	Active
supplier.tbl.5.csv	Hot (inferred)		2-1-2024 16:13	Block Blob	text/plain	138,24 MiB	Active
supplier.tbl.2.csv	Hot (inferred)		2-1-2024 15:44	Block Blob	text/plain	138,25 MiB	Active
supplier.tbl.3.csv	Hot (inferred)		2-1-2024 16:02	Block Blob	text/plain	138,29 MiB	Active
supplier.tbl.4.csv	Hot (inferred)		2-1-2024 16:35	Block Blob	text/plain	138,29 MiB	Active
part.tbl.1.csv	Hot (inferred)		2-1-2024 16:26	Block Blob	text/plain	2,30 GiB	Active
part.tbl.5.csv	Hot (inferred)		2-1-2024 15:45	Block Blob	text/plain	2,31 GiB	Active
part.tbl.4.csv	Hot (inferred)		2-1-2024 16:15	Block Blob	text/plain	2,31 GiB	Active
part.tbl.2.csv	Hot (inferred)		2-1-2024 16:34	Block Blob	text/plain	2,31 GiB	Active
part.tbl.3.csv	Hot (inferred)		2-1-2024 15:08	Block Blob	text/plain	2,31 GiB	Active
customer.tbl.3.csv	Hot (inferred)		2-1-2024 16:38	Block Blob	text/plain	2,32 GiB	Active
customer.tbl.5.csv	Hot (inferred)		2-1-2024 15:33	Block Blob	text/plain	2,32 GiB	Active
customer.tbl.2.csv	Hot (inferred)		2-1-2024 16:19	Block Blob	text/plain	2,32 GiB	Active
customer.tbl.4.csv	Hot (inferred)		2-1-2024 16:35	Block Blob	text/plain	2,32 GiB	Active
partsupp.tbl.1.csv	Hot (inferred)		2-1-2024 15:50	Block Blob	text/plain	11,51 GiB	Active
partsupp.tbl.3.csv	Hot (inferred)		2-1-2024 15:33	Block Blob	text/plain	11,55 GiB	Active
partsupp.tbl.5.csv	Hot (inferred)		2-1-2024 15:44	Block Blob	text/plain	11,55 GiB	Active
partsupp.tbl.2.csv	Hot (inferred)		2-1-2024 16:15	Block Blob	text/plain	11,55 GiB	Active
partsupp.tbl.4.csv	Hot (inferred)		2-1-2024 16:39	Block Blob	text/plain	11,55 GiB	Active
orders.tbl.1.csv	Hot (inferred)		2-1-2024 15:59	Block Blob	text/plain	16,79 GiB	Active
orders.tbl.2.csv	Hot (inferred)		2-1-2024 16:37	Block Blob	text/plain	16,87 GiB	Active
orders.tbl.4.csv	Hot (inferred)		2-1-2024 16:34	Block Blob	text/plain	16,96 GiB	Active
orders.tbl.3	Hot (inferred)		2-1-2024 16:18	Block Blob	text/plain	16,96 GiB	Active
orders.tbl.5.csv	Hot (inferred)		2-1-2024 16:01	Block Blob	text/plain	16,96 GiB	Active
lineitem.tbl.1	Hot (inferred)		2-1-2024 15:43	Block Blob	text/plain	75,42 GiB	Active
lineitem.tbl.1.csv	Hot (inferred)		2-1-2024 15:31	Block Blob	text/plain	75,42 GiB	Active
lineitem.tbl.2.csv	Hot (inferred)		2-1-2024 16:34	Block Blob	text/plain	75,70 GiB	Active
lineitem.tbl.3.csv	Hot (inferred)		2-1-2024 15:56	Block Blob	text/plain	76,07 GiB	Active
lineitem.tbl.5.csv	Hot (inferred)		2-1-2024 15:21	Block Blob	text/plain	76,08 GiB	Active
lineitem.tbl.4.csv	Hot (inferred)		2-1-2024 16:13	Block Blob	text/plain	76,08 GiB	Active

- TPC-H is free and contains nonsense data
- A little quirky to get running
- Very flexible to get the size you want

# Loading options



# Loading...

- Pipeline
- Notebook
- Dataflow Gen2



# Pipeline

Easy going



# Pipeline

Copy data

TPCH to Lakehouse

Source Destination Mapping Settings

Destination type: Workspace

Destination data store type: Lakehouse

File name: TCPH-RAW

Format: Binary

Settings

This screenshot shows a data pipeline configuration interface. At the top, there is a modal window titled "Copy data" containing a single item: "TPCH to Lakehouse". Below this, the main interface has tabs for "Source", "Destination", "Mapping", and "Settings", with "Destination" currently selected. Under "Destination", the "type" is set to "Workspace" (radio button selected) and the "data store type" is set to "Lakehouse" (dropdown menu). A dropdown menu for "File name" lists options: "Lakehouse", "KQL Database", "Warehouse", and "TCPH-RAW", with "TCPH-RAW" currently selected. Below the dropdown are buttons for "Refresh", "Open", and "New". At the bottom, there are fields for "Format" (set to "Binary") and "Settings".

# Pipeline

General	Source	Destination	Mapping	Settings
Intelligent throughput optimization <small>i</small>				<input type="button" value="Auto"/> <input type="button" value="▼"/>
				<input type="checkbox"/> Use custom value
Degree of copy parallelism <small>i</small>				<input type="button" value="Auto"/> <input type="button" value="▼"/>
Data consistency verification <small>i</small>				<input type="checkbox"/>
Fault tolerance <small>i</small>				<input type="button" value="Skip missing files"/> <input type="button" value="▼"/>
				<input type="checkbox"/>

# Comparison

## Pipeline F64

Throughput	335.235 MB/s
Total duration	00:33:14

### Duration breakdown

Start time 1/2/2024, 4:54:14 PM

Optimized throughput ⓘ Balanced

Used parallel copies ⓘ 638



## Pipeline F2

Throughput	341.603 MB/s
Total duration	00:32:39

### Duration breakdown

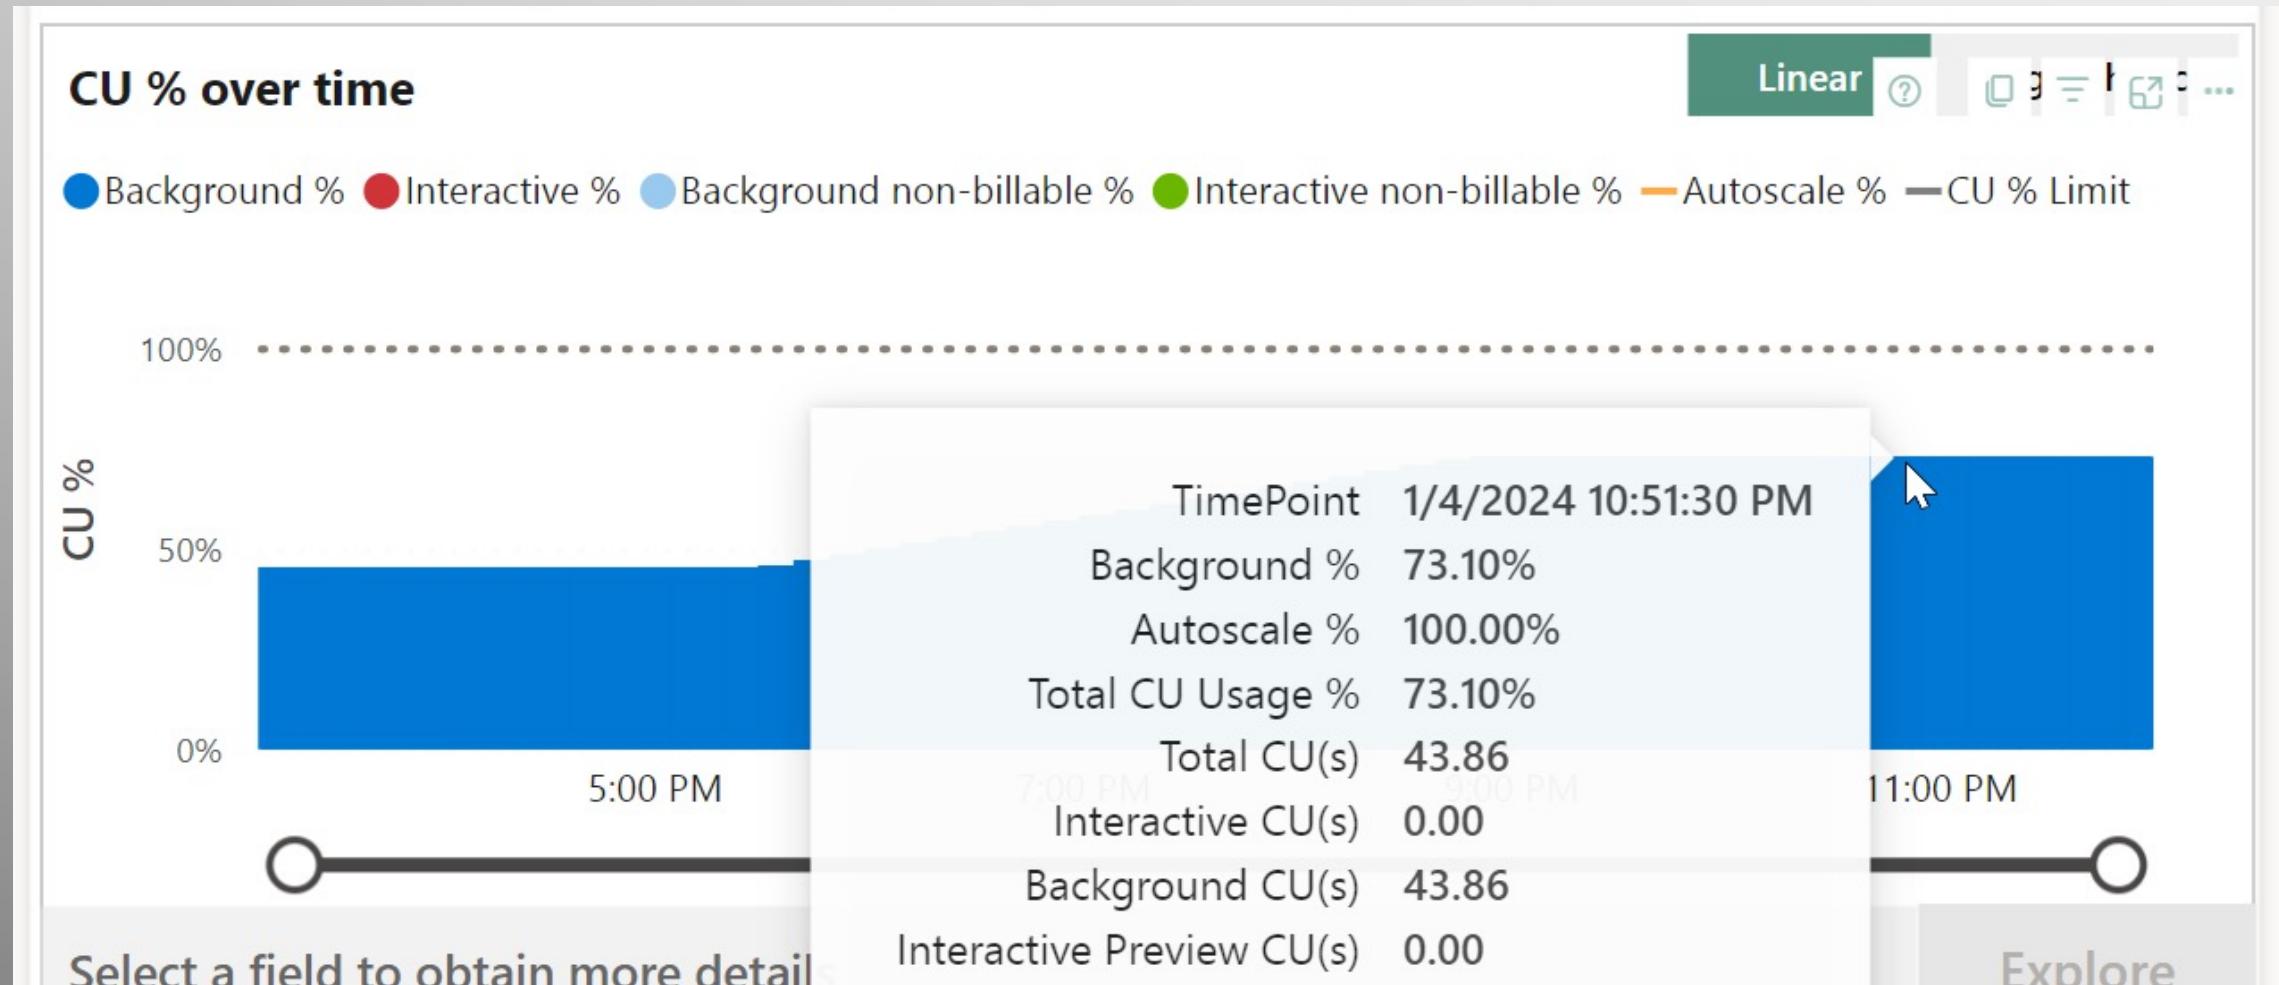
Start time 1/4/2024, 2:29:49 PM

Optimized throughput ⓘ Balanced

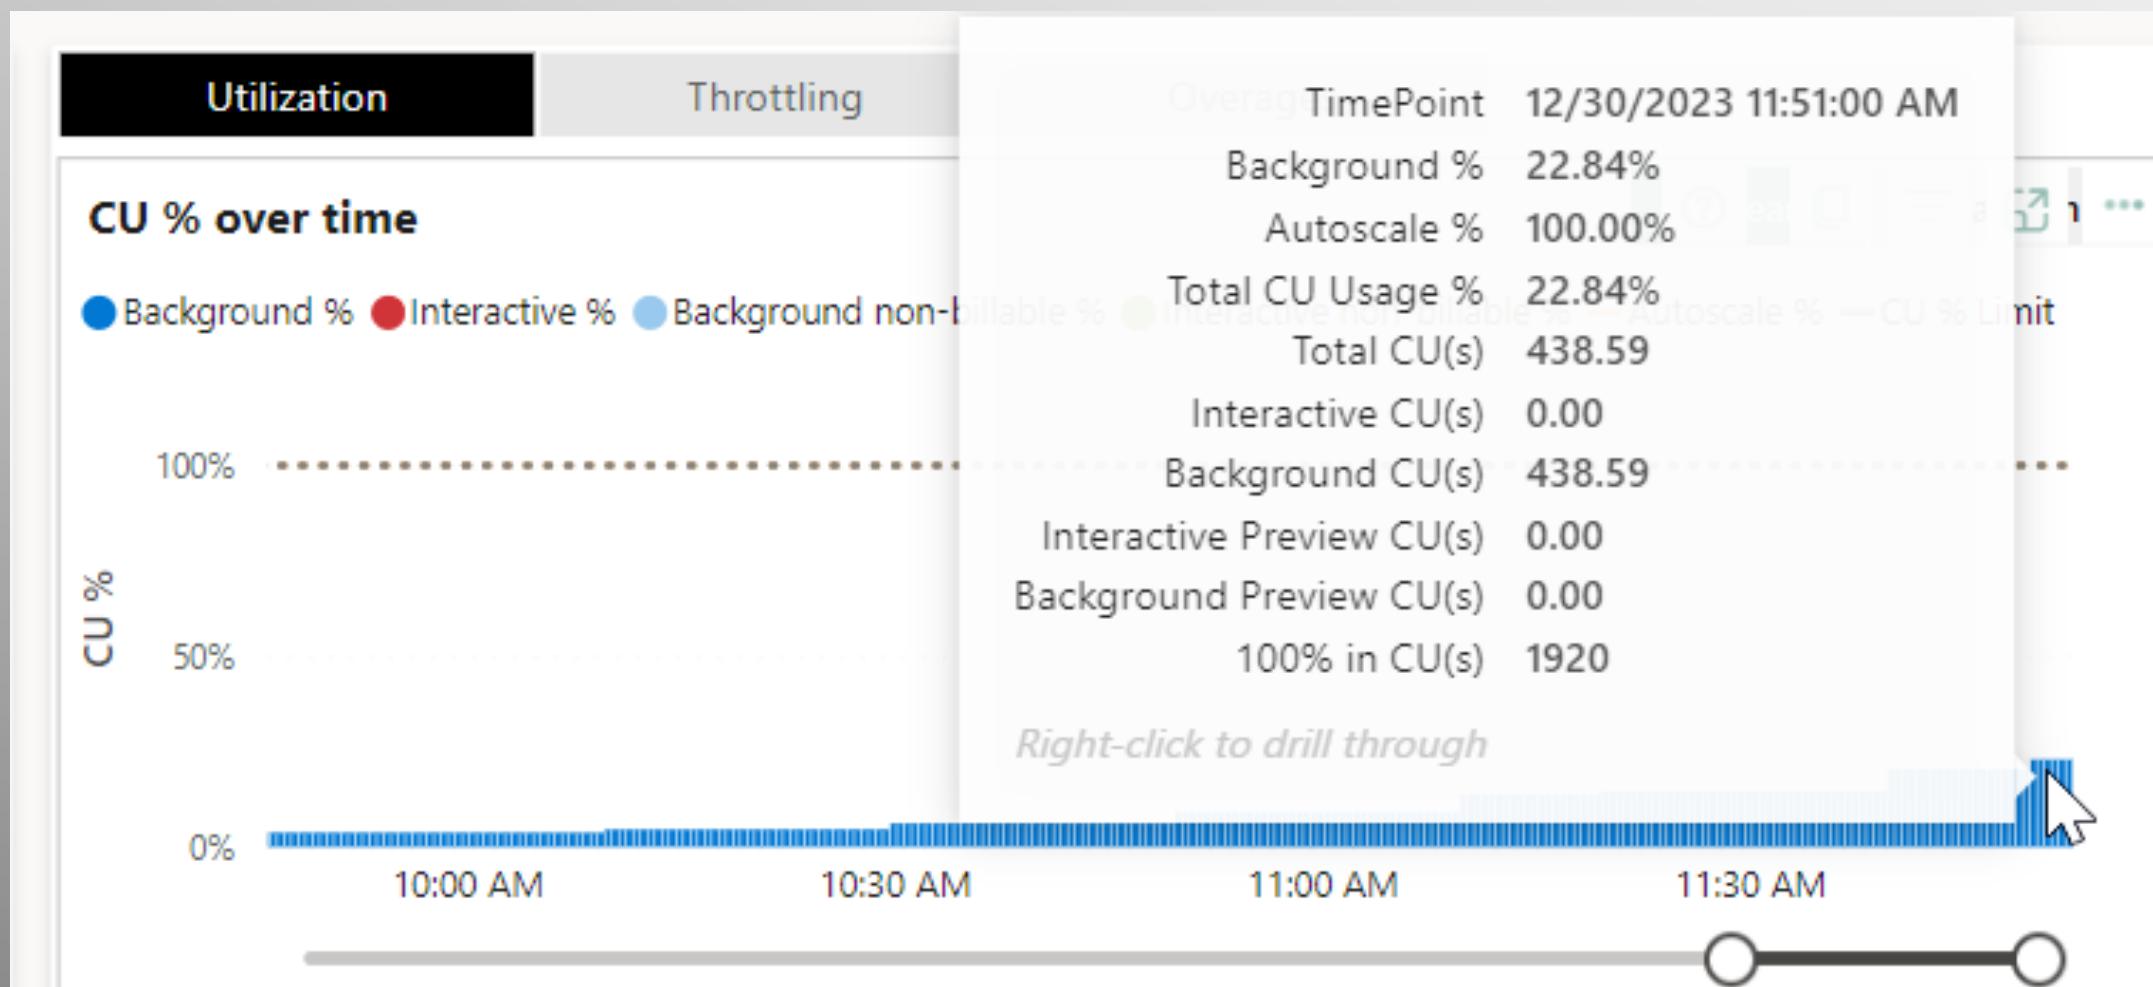
Used parallel copies ⓘ 638



# Pipeline F2



# Pipeline F64



# Pipeline F64

- **Timepoint CU(s)** - This metric is the total CU(s) divided by 2,880, and is used to determine how much CU(s) the background operation contributes to this timepoint.

Status	User	Duration (s)	Total CU (s)	Timepoint CU (s)	Throttling (s)	% of Base capacity	Billing type
Success	Power BI Service	920	318,240.0000	110.5000	0	5.76%	Billable
Success	Power BI Service	938	243,000.0000	84.3750	0	4.39%	Billable
Success	Power BI Service	459	145,080.0000	50.3750	0	2.62%	Billable
Success	Power BI Service	701	120,960.0000	42.0000	0	2.19%	Billable
Success	Power BI Service	1,087	88,560.0000	30.7500	0	1.60%	Billable
Success	Power BI Service	218	50,040.0000	17.3750	0	0.90%	Billable
Success	Power BI Service	179	46,800.0000	16.2500	0	0.85%	Billable
Success	Power BI Service	188	46,440.0000	16.1250	0	0.84%	Billable

$$(24 * 60 * 60) / 2880 = 30 \text{ (seconds)}$$

# Pipeline

F2 uses 73% capacity

F64 uses 2,2% capacity

What else is going on?

# Notebook

Let's code!



# Notebook (CSV to Parquet)

```
1  folderPath = "Files/TCPH-RAW/files/customer.tbl.*.csv"
2  table_name = "raw_customers"
3
4  schema = StructType([
5      StructField("C_CUSTKEY", IntegerType()),
6      StructField("C_NAME", StringType()),
7      StructField("C_ADDRESS", StringType()),
8      StructField("C_NATIONKEY", IntegerType()),
9      StructField("C_PHONE", StringType()),
10     StructField("C_ACCTBAL", DoubleType()),
11     StructField("C_MKTSEGMENT", StringType()),
12     StructField("C_COMMENT", StringType()),
13     StructField("C_EEmpty", StringType())
14 ])
15
16 df = spark.read.csv(
17     folderPath,
18     header=False,
19     mode="DROPMALFORMED",
20     sep="|",
21     schema=schema
22 )
23 df.write.mode("overwrite").format("delta").save("Tables/" + table_name)
```

# Notebook (CSV to Parquet) F2

✓ 2 min 29 sec -Command executed in 2 min 28 sec 902 ms by Admin on 11:15:34 AM, 1/05/24

Tasks	Duration	Processed	Data read	Data written
2/2 succeeded	198 ms	19 rows	12.12 KB	0 B
76/76 succeeded	1 min 4 sec 806 ms	120,000,000 rows	0 B	6.3 GB
5/5 succeeded	1 min 17 sec 26 ms	120,000,000 rows	6.3 GB	4.68 GB

# Notebook (CSV to Parquet) F64

✓ 1 min 13 sec -Command executed in 1 min 13 sec 264 ms by Admin on 1:46:41 PM, 1/05/24

Tasks	Duration	Processed	Data read	Data written
1/1 succeeded	186 ms	16 rows	5.96 KB	5.25 KB
50/50 succeeded	466 ms	58 rows	5.25 KB	4.37 KB
1/1 succeeded	47 ms	50 rows	4.37 KB	0 B
76/76 succeeded	33 sec 632 ms	120,000,000 rows	0 B	6.3 GB
5/5 succeeded	35 sec 49 ms	120,000,000 rows	6.3 GB	4.68 GB

# Notebook (CSV to Parquet)

```
1  folderPath = "Files/TCPH-RAW/files/lineitem.tbl.*.csv"
2  table_name = "raw_lineitems"
3
4  schema = StructType([
5      StructField("L_ORDERKEY", IntegerType()),
6      StructField("L_PARTKEY", IntegerType()),
7      StructField("L_SUPPKEY", IntegerType()),
8      StructField("L_LINENUMBER", IntegerType()),
9      StructField("L_QUANTITY", IntegerType()),
10     StructField("L_EXTENDEDPRICE", DoubleType()),
11     StructField("L_DISCOUNT", DoubleType()),
12     StructField("L_TAX", DoubleType()),
13     StructField("L_RETURNFLAG", StringType()),
14     StructField("L_LINESSTATUS", StringType()),
15     StructField("L_SHIPDATE", StringType()),
16     StructField("L_COMMITDATE", StringType()),
17     StructField("L_RECEIPTDATE", StringType()),
18     StructField("L_SHIPINSTRUCT", StringType()),
19     StructField("L_SHIPMODE", StringType()),
20     StructField("L_COMMENT", StringType()),
21     StructField("L_Empty", StringType())
22 ])
23
24 df = spark.read.csv(
25     folderPath,
26     header=False,
27     mode="DROPMALFORMED",
28     sep="|",
29     schema=schema
30 )
31 df.write.mode("overwrite").format("delta").save("Tables/" + table_name)
```

# Notebook (CSV to Parquet) F2

[3] ✓ 2 h 4 min -Command executed in 2 h 4 min 55 sec 205 ms by Admin on 1:20:30 PM, 1/05/24

▼ \*  Spark jobs (11 of 11 succeeded) \*  Log

3036/3044 succeeded, 8 failed	1 h 17 min 29 sec 991 ms	4,294,965,612 rows	6 GB	174.67 GB
136/136 succeeded	47 min 19 sec 41 ms	4,294,965,612 rows	174.67 GB	77.44 GB

# Notebook (CSV to Parquet) F64

[62]

✓ -Command executed in 7 min 46 sec 31 ms by Admin on 7:38:02 PM, 1/02/24

---

3036/3044 succeeded, 8 failed

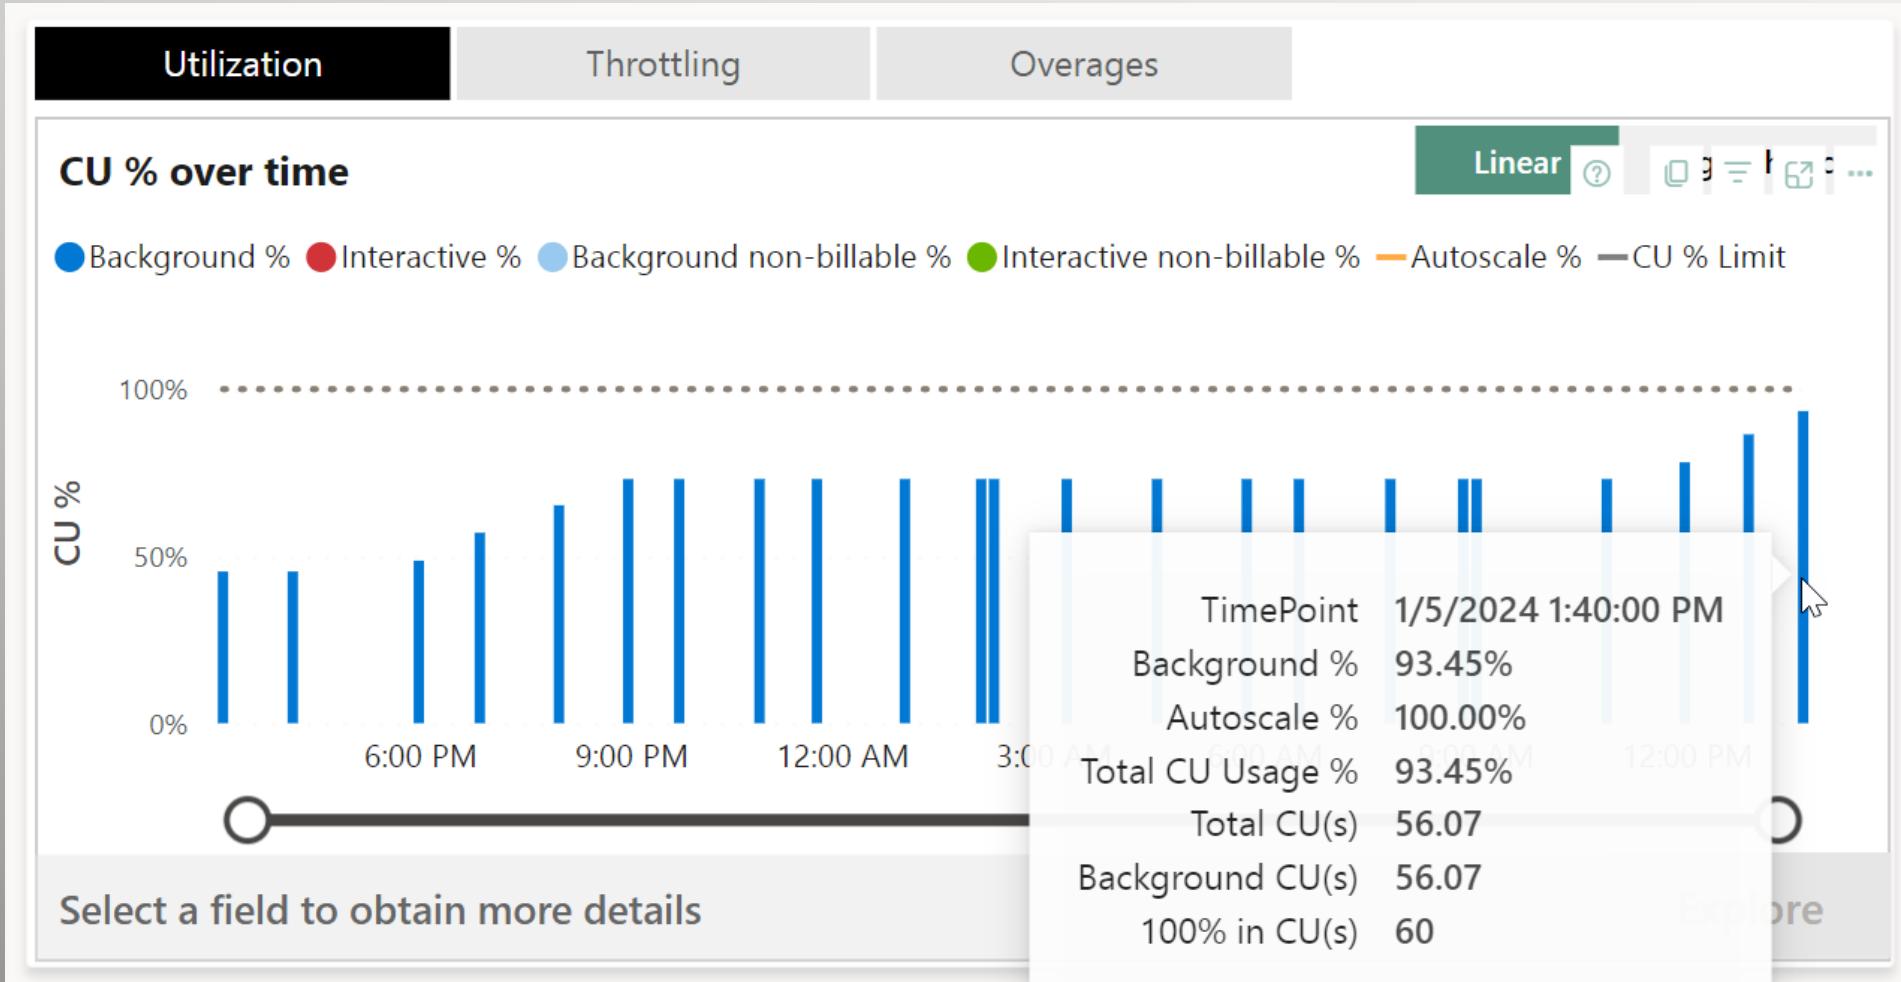
---

1 h 17 min 29 sec 991 ms    4,294,965,612 rows    6 GB    174.67 GB

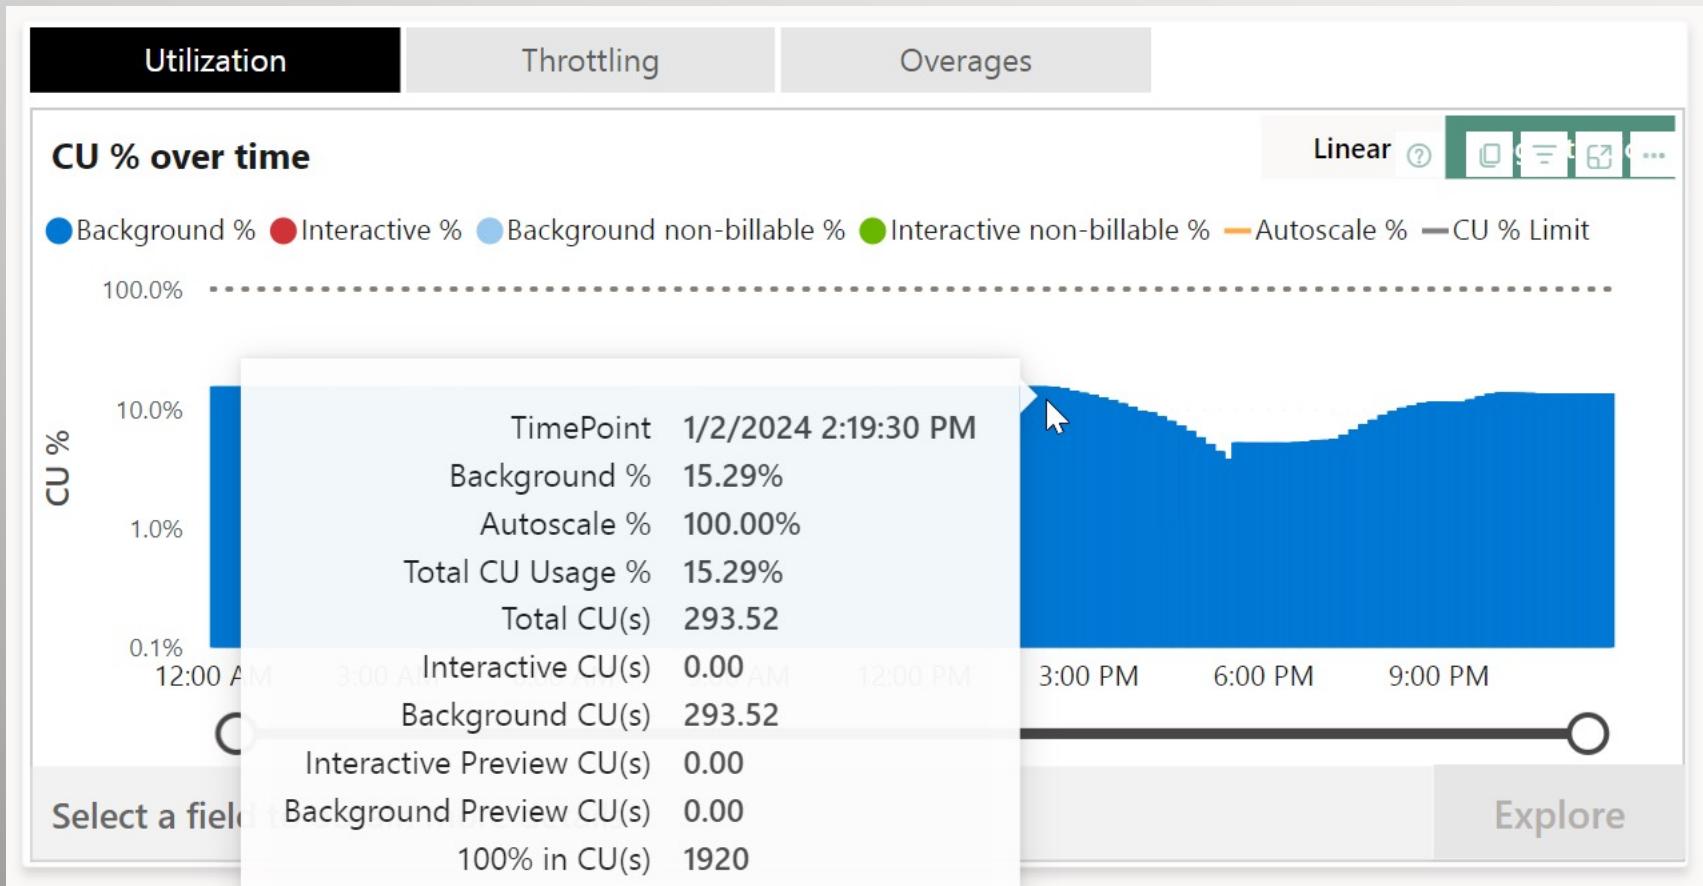
136/136 succeeded

47 min 19 sec 41 ms    4,294,965,612 rows    174.67 GB    77.44 GB

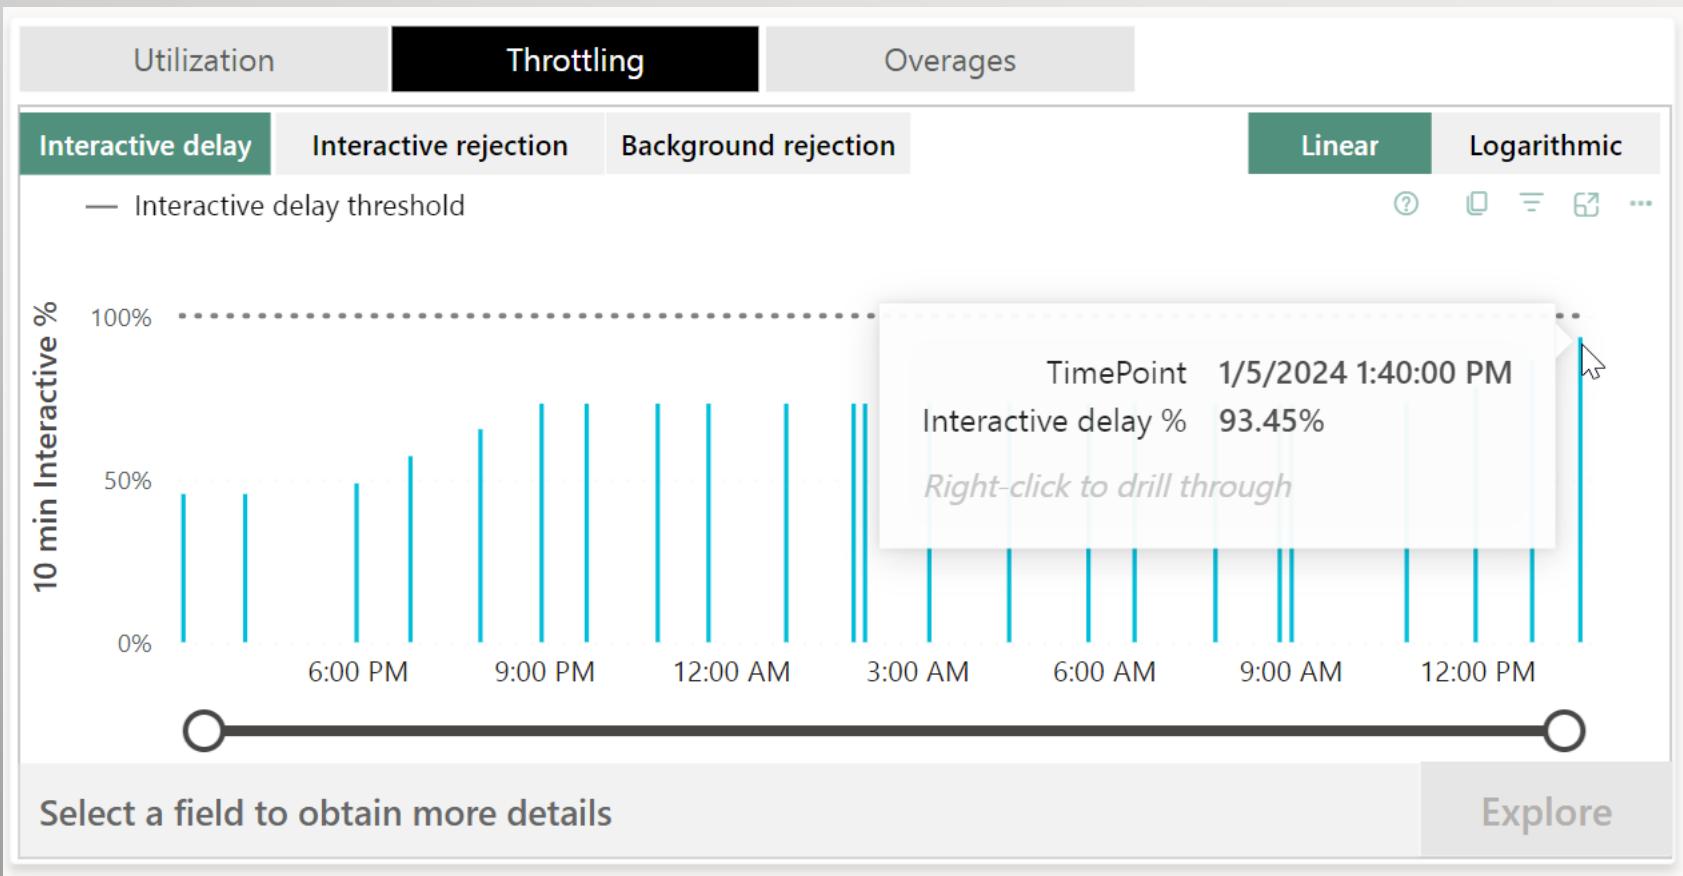
# Notebook (CSV to Parquet) Metrics



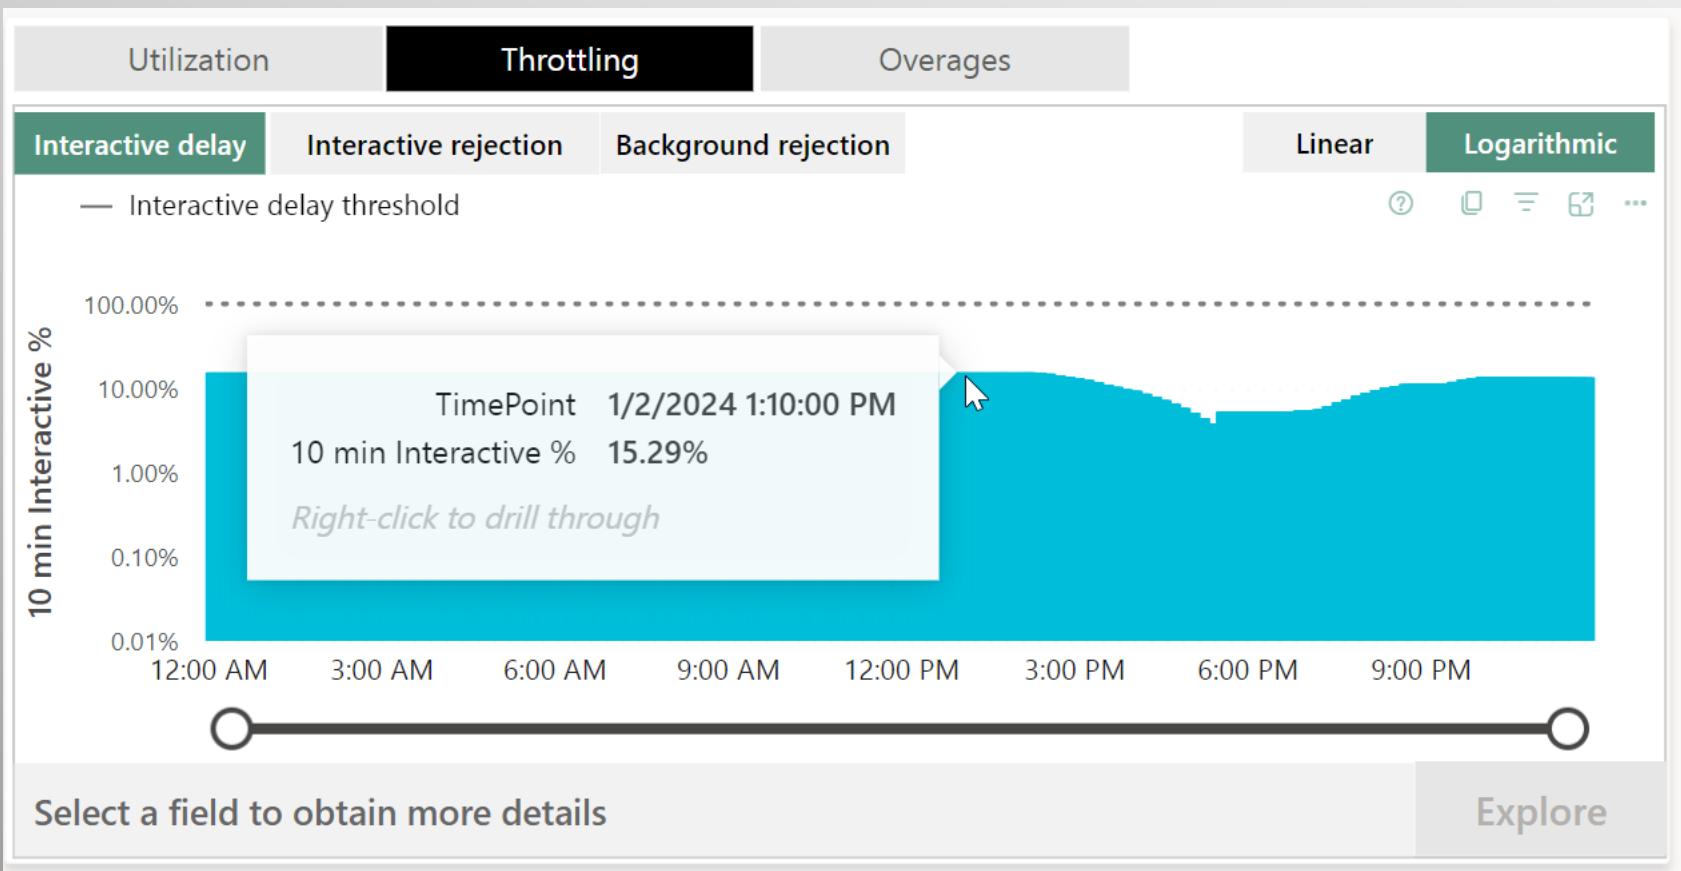
# Notebook (CSV to Parquet) Metrics



# Notebook (CSV to Parquet) Metrics



# Notebook (CSV to Parquet) Metrics



# Create a table with some extra columns F2

```
1  from pyspark.sql.functions import col, year, month, quarter
2
3  table_name = 'fact_orders_2cu'
4
5  df = spark.sql("SELECT * FROM Demo_Lakehouse_2CU.raw_orders")
6  df = df.withColumn('Year', year(col("o_orderdate")))
7  df = df.withColumn('Quarter', quarter(col("o_orderdate")))
8  df = df.withColumn('Month', month(col("o_orderdate")))
9
10 df.write.mode("overwrite").format("delta").partitionBy("Year", "Quarter").save("Tables/" + table_name)
```

✓ 18 min 7 sec -Command executed in 18 min 7 sec 97 ms by Admin on 2:56:34 PM, 1/05/24

s	773,741,822 rows	16.28 GB	23.73 GB
s	773,741,822 rows	23.73 GB	15.85 GB

# Create a table with some extra columns F64

```
1 from pyspark.sql.functions import col, year, month, quarter
2
3 table_name = 'fact_orders'
4
5 df = spark.sql("SELECT * FROM FabricLakehouse.raw_orders")
6 df = df.withColumn('Year', year(col("o_orderdate")))
7 df = df.withColumn('Quarter', quarter(col("o_orderdate")))
8 df = df.withColumn('Month', month(col("o_orderdate")))
9
10 df.write.mode("overwrite").format("delta").partitionBy("Year", "Quarter").save("Tables/" + table_name)
```

✓ 8 min 21 sec -Command executed in 8 min 20 sec 680 ms by Admin on 12:00:16 PM, 1/07/24

s	773,741,822 rows	16.28 GB	23.73 GB
s	773,741,822 rows	23.73 GB	15.85 GB

# Join two huge tables F2

```
1  from pyspark.sql.functions import col, year, month, quarter
2
3  table_name = 'fact_orderlines'
4
5  dforder = spark.sql("SELECT * FROM Demo_Lakehouse_2CU.raw_orders")
6  dforder = dforder.withColumn('Year', year(col("o_orderdate")))
7  dforder = dforder.withColumn('Quarter', quarter(col("o_orderdate")))
8  dforder = dforder.withColumn('Month', month(col("o_orderdate")))
9
10 dfline = spark.sql("SELECT * FROM Demo_Lakehouse_2CU.raw_lineitems")
11
12 dforderline = dforder.join(dfline,dforder.O_ORDERKEY == dfline.L_ORDERKEY, how='Inner')
13
14 dforderline.write.mode("overwrite").format("delta").partitionBy("Year","Quarter").save("Tables/" + table_name)
```

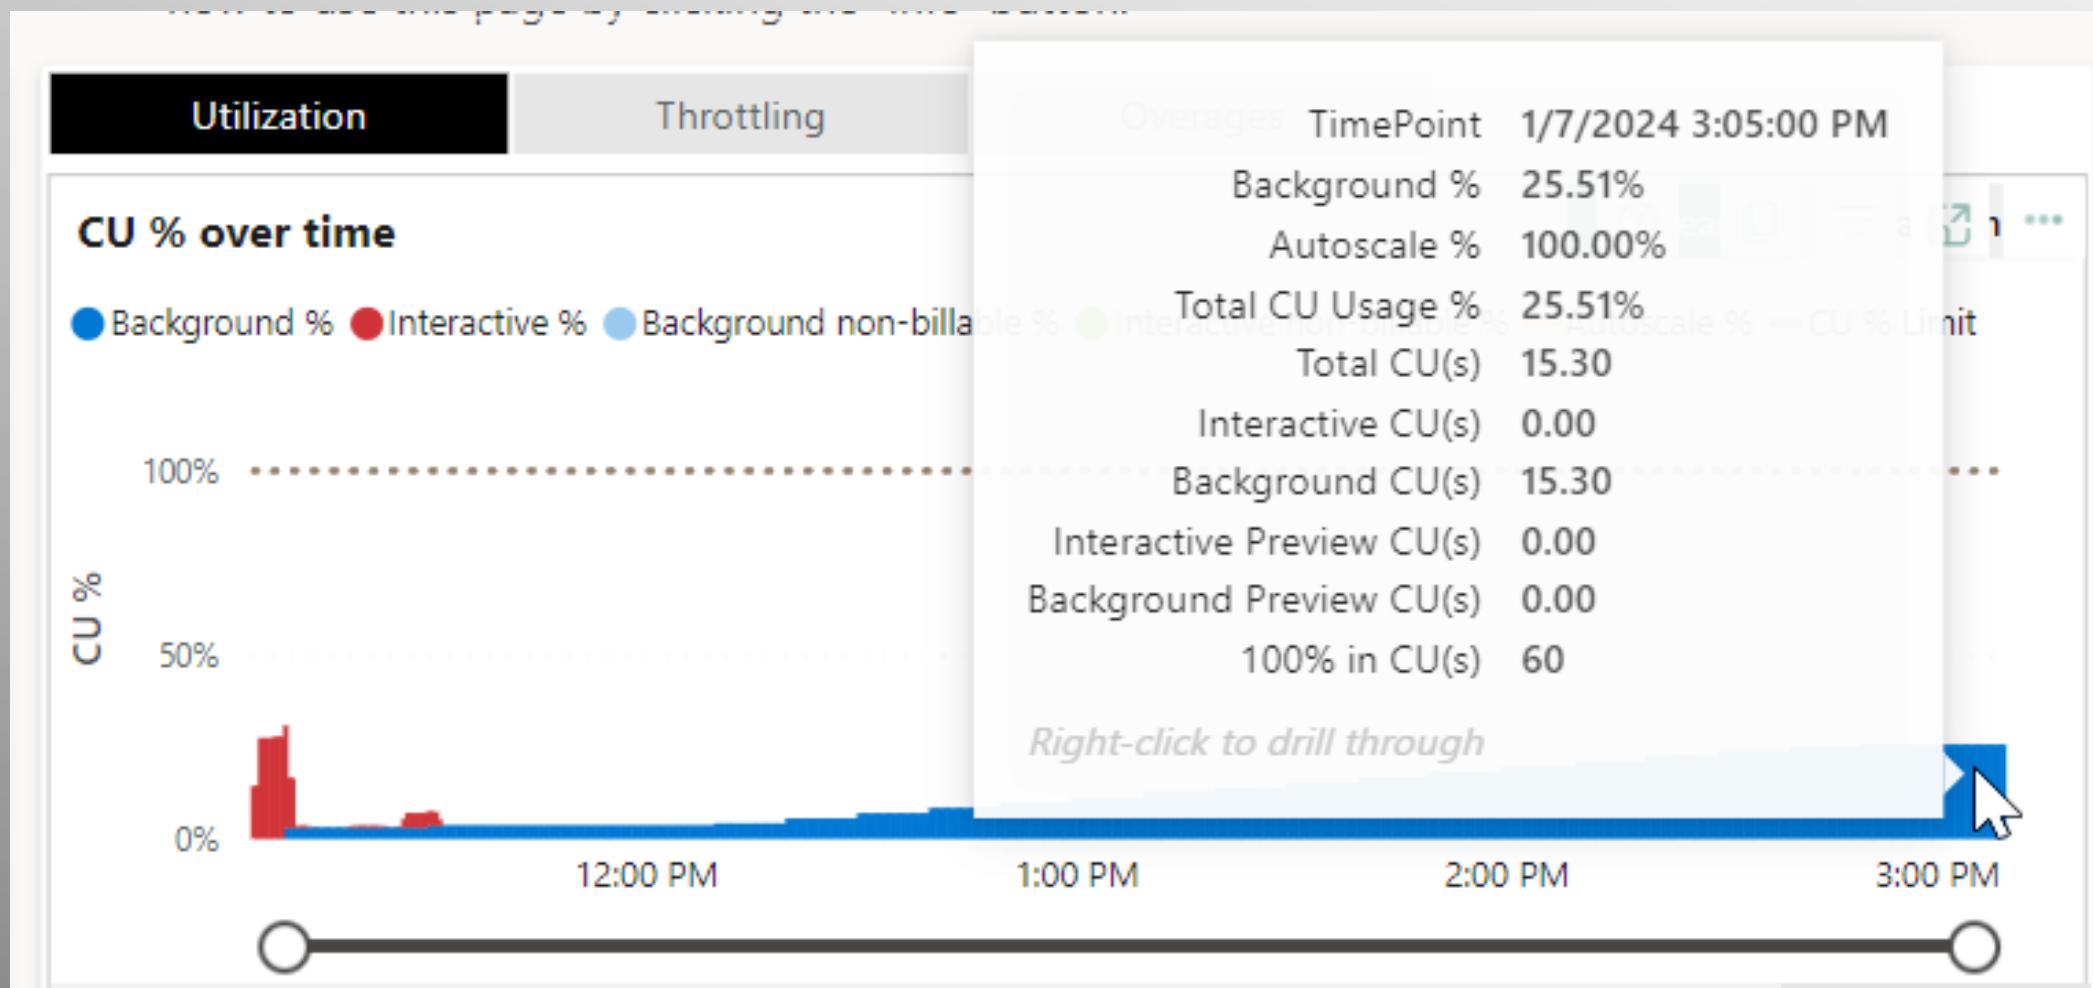
✓ 2 h 16 min -Command executed in 2 h 16 min 23 sec 908 ms by Admin on 2:22:06 PM, 1/07/24

# Join two huge tables F64

```
1  from pyspark.sql.functions import col, year, month, quarter
2
3  table_name = 'fact_orderlines'
4
5  dforder = spark.sql("SELECT * FROM FabricLakehouse.raw_orders")
6  dforder = dforder.withColumn('Year', year(col("o_orderdate")))
7  dforder = dforder.withColumn('Quarter', quarter(col("o_orderdate")))
8  dforder = dforder.withColumn('Month', month(col("o_orderdate")))
9
10 dfline = spark.sql("SELECT * FROM FabricLakehouse.raw_lineitems")
11
12 dforderline = dforder.join(dfline,dforder.O_ORDERKEY == dfline.L_ORDERKEY, how='Inner')
13
14 dforderline.write.mode("overwrite").format("delta").partitionBy("Year","Quarter").save("Tables/" + table_name)
```

✓ -Apache Spark session ready in 2 min 39 sec 346 ms. Command executed in 14 min 51 sec 480 ms by Admin on 9:33:17 PM, 1/02/24

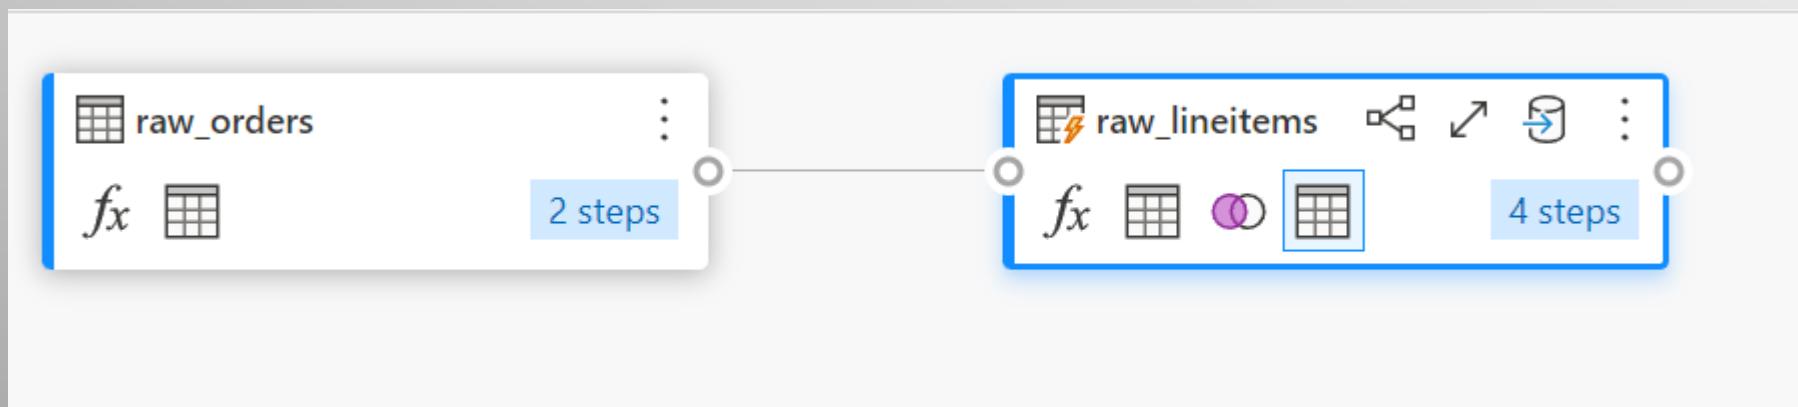
# Join CU requests



# Pipeline Gen2

Let's ....





▼ Applied steps

 Source

 Navigation 1

 Merged queries

  Navigation



ⓘ DataSource.Error: Microsoft SQL: The SQL pool is warming up. Please try again.

Details

Reason = DataSource.Error

DataSourceKind = Lakehouse

DataSourcePath = Lakehouse

Message = The SQL pool is warming up. Please try again.

ErrorCode = -2146232060

Number = 42109

Class = 20

State = 1

 Failed	01:56:25	On demand
--	----------	-----------

An external error occurred while refreshing the dataflow: Unable to refresh the dataflow because your organization's Fabric compute capacity has exceeded its limits. Try again later. Learn more at  
<https://aka.ms/capacitymessages>

TimePoint 3/7/2024 7:38:30 AM  
Add % 16.76%  
Add % 16.76%  
12:00 AM Burndown % 0.00% 6:00 AM  
Cumulative % (Right axis) 22102.83%  
Expected minutes to burndown 110.51

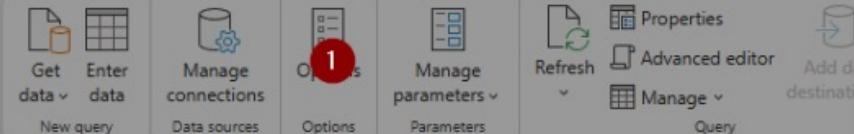
TimePoint	3/6/2024 10:09:30 PM
Background %	114.52%
Autoscale %	100.00%
Total CU Usage %	114.52%
Total CU(s)	68.71
Interactive CU(s)	0.00
Background CU(s)	68.71
Interactive Preview CU(s)	0.00
Background Preview CU(s)	0.00
100% in CU(s)	60

# Pipeline Gen2 Turbo button

Let's!



Home Transform Add column View Help



Queries [2]

- raw\_lineitems**
- raw\_orders

	1 <sup>2</sup> 3 L_ORDERKEY	1 <sup>2</sup> 3 L_PARTKEY	1 <sup>2</sup> 3 L_SUPPKEY	1 <sup>2</sup> 3 L
1	19710342	25955423	955424	
2	19746561	25696657	1946663	
3	19834695	99582159	3332217	
4	19895271	45952653	2202663	
5	19969987	6655797	2905799	
6	20030247	61634116	4134141	
7	20087776	41856442	3106451	
8	19708583	85231518	2731553	
9	19745671	94493460	1993497	
10	19834242	24735293	2235302	
11	19886213	45339808	2839827	
12	19969573	55082390	3832424	
13	20028833	39330177	4330178	
14	20087685	44762394	1012403	

**Options****Global**

Allow use of fast copy connectors  3

General

Data load

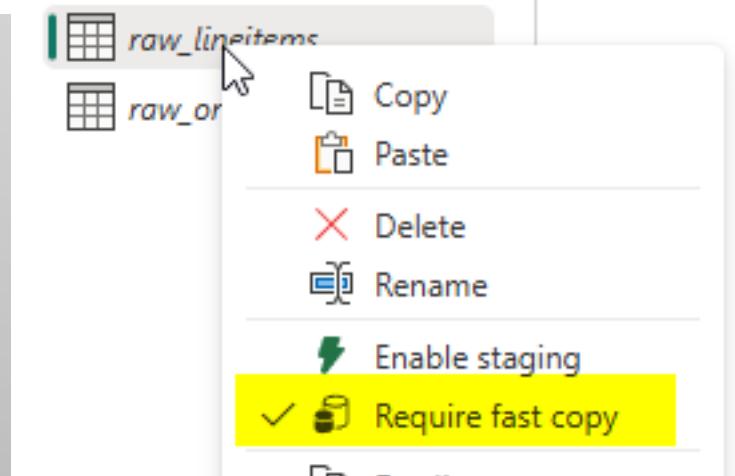
Diagnostics

**Dataflow**

Data load

Regional settings

Privacy

**Scale** 2**Queries [2]**

# Comparison

## Without Fast Copy

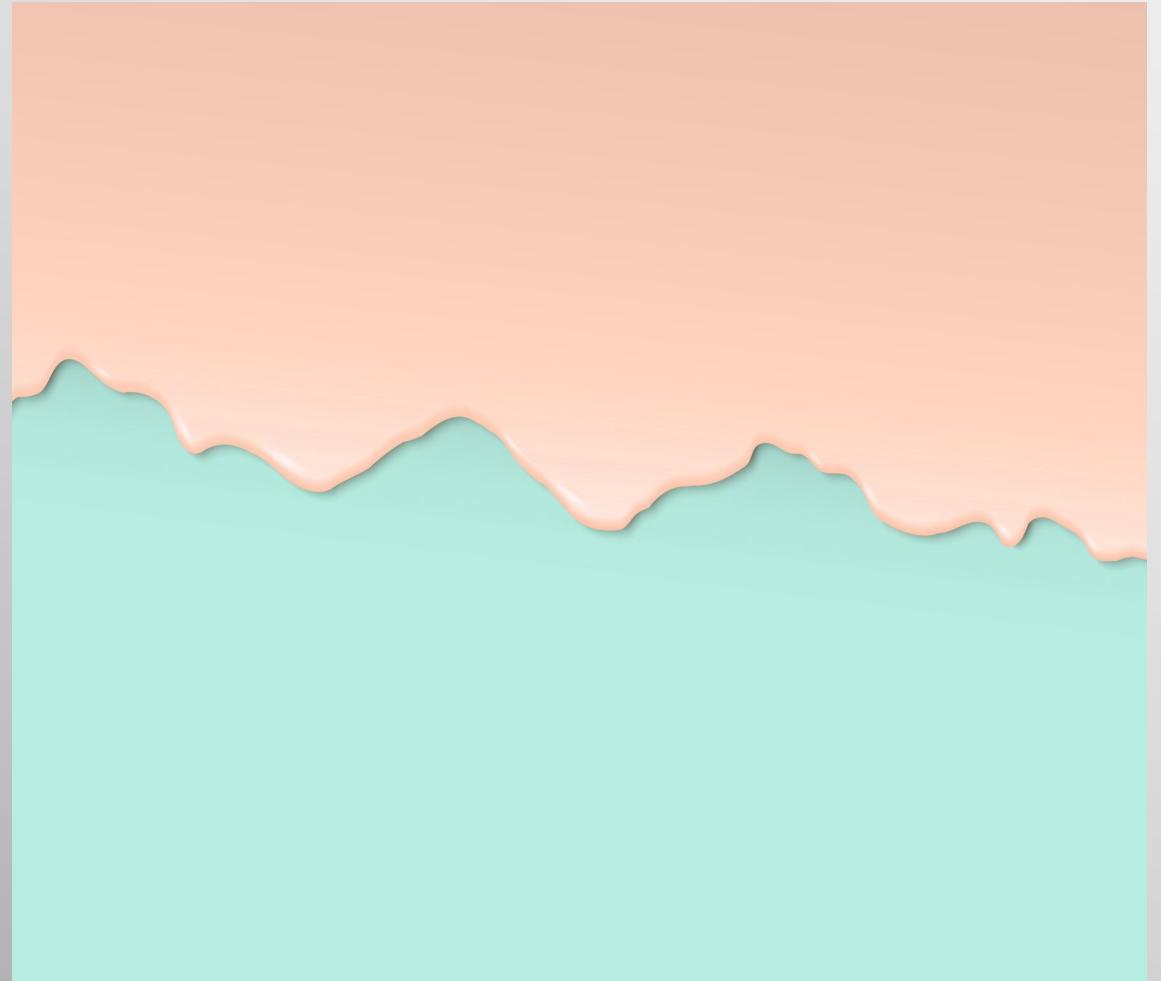
- Ingestion fails
- F2 and F64 fail
- 394 CU's
- Joining tables fails

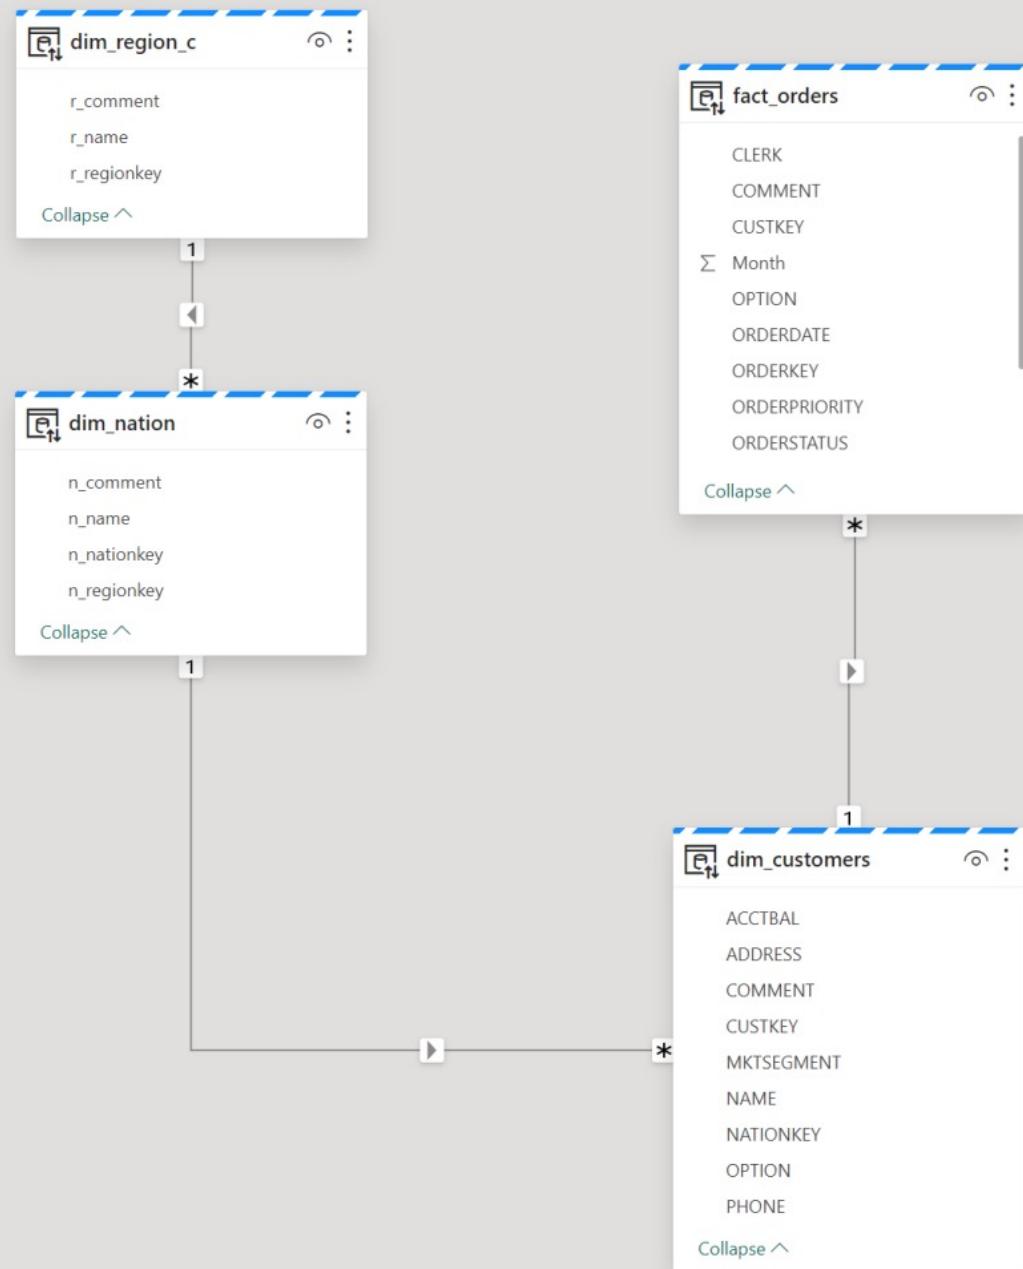
## Fast Copy

- Ingestion succeeds
- F64 succeeds
- 78 CU's
- Joining tables fails

# PowerBI

Are you serious?





## Couldn't load the data for this visual

X

The resultset of a query to external data source has exceeded the maximum allowed size of '1000000' rows.

Please try again later or contact support. If you contact support, please provide these details.

See details ▾

[Contact support](#)

[Close](#)



Unable to load model due to reaching capacity limits [See details](#)

## Unable to load model due to reaching capacity limits

X

Unable to open this report because your organization's compute capacity has exceeded its limits. Try again later.

Please try again later or contact support. If you contact support, please provide these details.

See details ▾

[Learn more about solving this issue](#)

[Close](#)

## Overages % over time

?

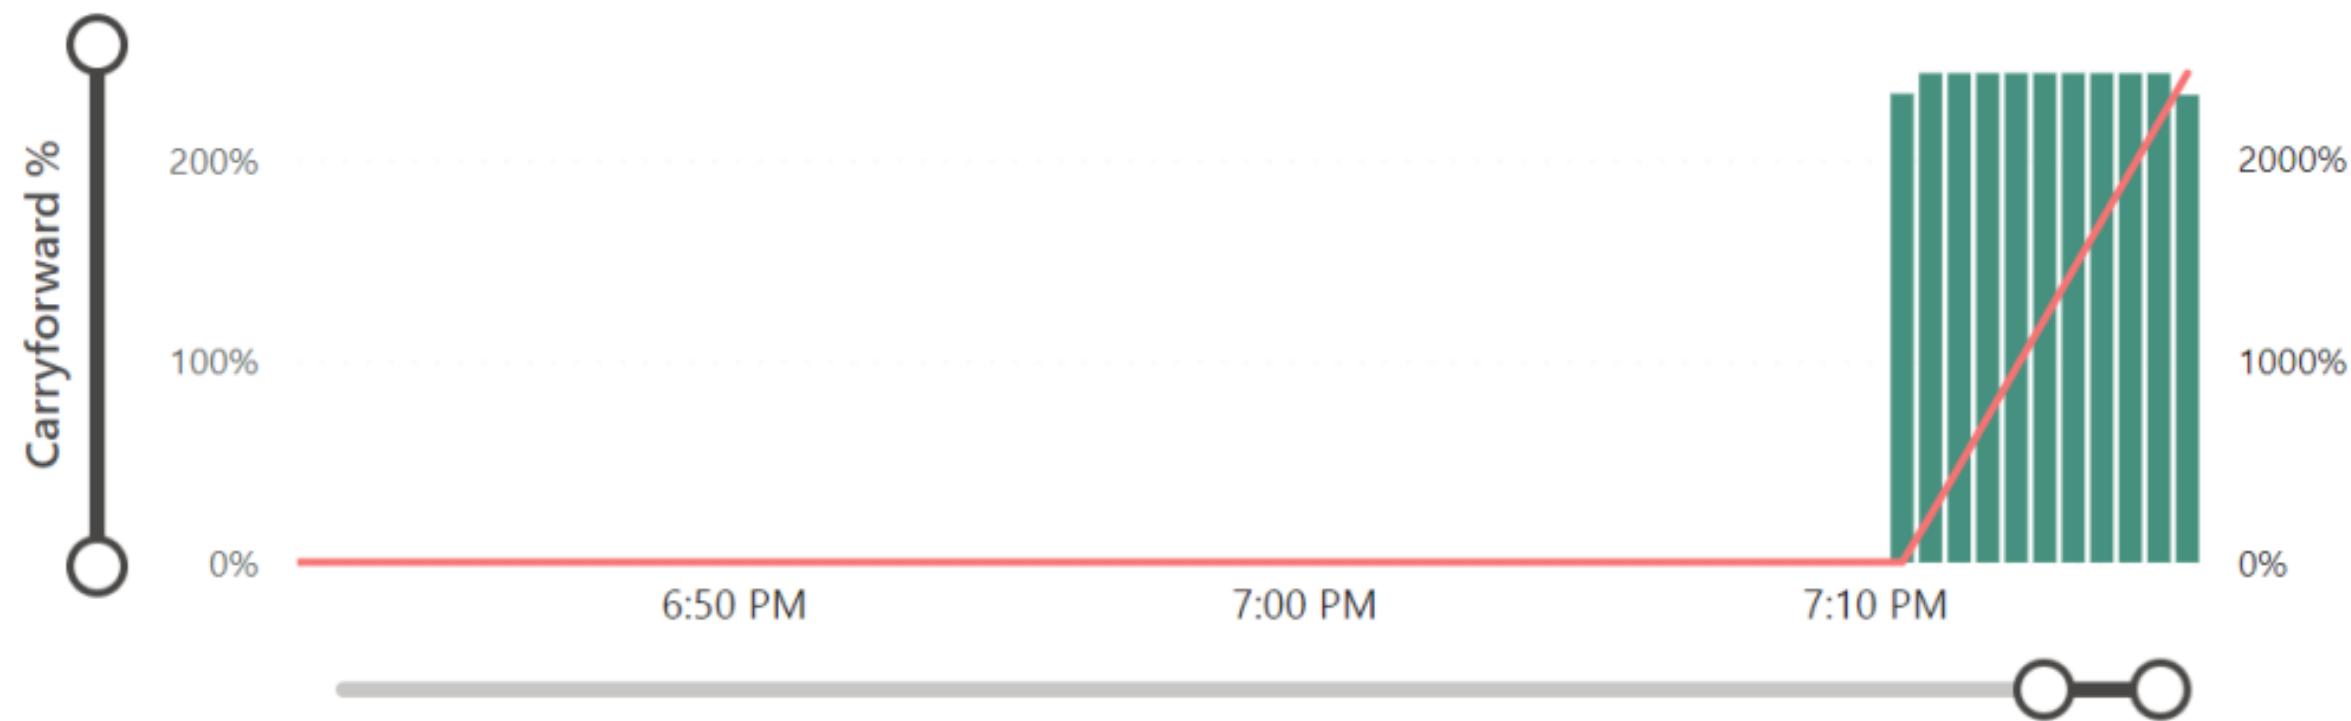
□

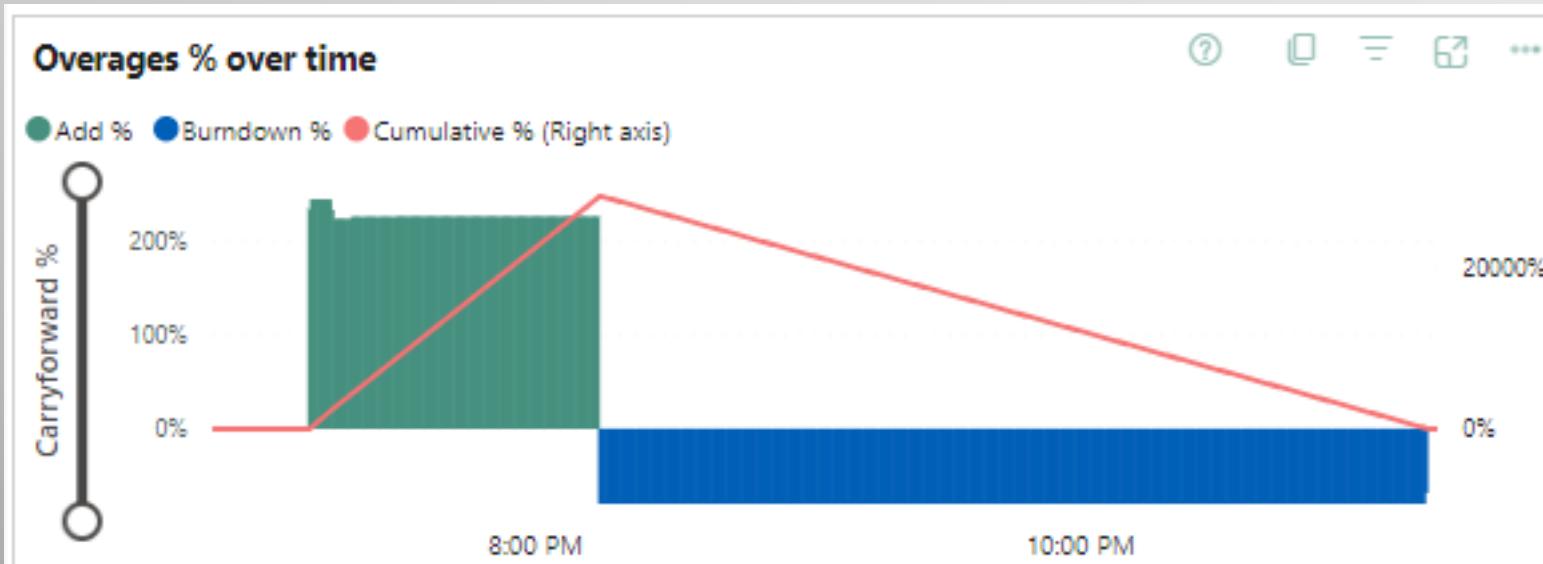
≡

↔

...

● Add % ● Burndown % ● Cumulative % (Right axis)





# Faster burndown

Can I have some more CU's please?



	TimePoint	2/28/2024 9:43:30 PM
9:35 PM	Burndown %	-42.48%
	Add %	0.00%
	Burndown %	-42.48%
tain more	Cumulative % (Right axis)	1406.49%
more	Expected minutes to burndown	7.03

9:35 PM	TimePoint	2/28/2024 9:44:00 PM
	Burndown %	-85.62%
	Add %	0.00%
	Burndown %	-85.62%
more	Cumulative % (Right axis)	341.00%
more	Expected minutes to burndown	1.71

# Summaries



# Notebook Metrics overview

TableName	Rows	F2	F64	Factor	CU(s)
Pipeline Ingestion	n/a	32 minutes 39 seconds	33 minutes 14 seconds	1x	43
Customers	120 million	2 minutes 29 seconds	1 minute 13 seconds	2x	56
LineItems	4.2 billion	2 hours, 4 minutes 55 seconds	7 minutes 46 seconds	17x	82
Fact_orders	773 million	18 minutes	8 minutes	2.25x	26
Order_lines	4.2 billion	2 hours 16 minutes	14 minutes 51 seconds	9x	15
<b>Sum</b>		<b>~5 hours</b>	<b>~65 minutes</b>	<b>5x</b>	<b>222</b>

# Fabric licenses

SKU	Capacity Units (CU)	Power BI SKU	Power BI v-cores	Spark Cores
F2	2	-	0.25	4
F4	4	-	0.5	8
F8	8	EM/A1	1	16
F16	16	EM2/A2	2	32
F32	32	EM3/A3	4	64
F64	64	P1/A4	8	128

# Results

- [Install the Microsoft Fabric capacity metrics app - Microsoft Fabric | Microsoft Learn](#)
- Test your workload as the CU(s) are like Sql Server Query costs
- Compare numbers and review duration and CU(s)

# Resources

- TPC-H
  - [https://www\(tpc.org/tpch/](https://www(tpc.org/tpch/)
  - <https://github.com/dimitri/tpch-citus/blob/master/schema/tpch-schema.sql>
  - <https://docs.deistercloud.com/content/Databases.30/TPCH%20Benchmark.90/index.xml?embedded=true>
  - [https://support.hpe.com/hpsc/public/docDisplay?docId=sf000078704en\\_us&docLocale=en\\_US](https://support.hpe.com/hpsc/public/docDisplay?docId=sf000078704en_us&docLocale=en_US)

# Resources

- Fabric Capacity App
  - <https://gennakerdata.com/2023/12/12/capacity-metrics-app-analysis-english-version/>
  - <https://learn.microsoft.com/en-us/fabric/enterprise/metrics-app-install?tabs=1st>
  - <https://blog.fabric.microsoft.com/en-us/blog/capacity-platform-updates-for-pause-resume-and-capacity-metrics-for-copilot-and-vnet-gateways/>



**Thank you for attending!**

**Any questions?**